

Producer/consumer based on a FIFO Queue

```
public produce(Object x) {  
    mutex.lock();  
    try {  
        queue.enq(x);  
    } finally {  
        mutex.unlock();  
    }  
}
```

The Need for Modular Synchronization

Suppose queue is bounded:

- enq may block until queue has room
- decision whether to block depends on internal state of the queue

Multiple producers/consumers:

- every thread needs to keep track of the lock, the queue state, etc.

The Need for Modular Synchronization

Suppose queue is bounded:

- enq may block until queue has room
- decision whether to block depends on internal state of the queue

Multiple producers/consumers:

- every thread needs to keep track of the lock, the queue state, etc.

not scalable

Modular Synchronization

Let queue handle its own synchronization

- queue has its own lock
 - acquired by each method call
 - released when the call returns
- if thread enqueues on a full queue
 - queue itself detects the problem
 - suspend the caller and resume when the queue has room

Conditions

- a condition object is associated with a lock
- condition objects allow a thread to
 - temporarily release the lock and suspend itself until awoken by another thread
 - awake other threads that are currently suspended waiting for that condition

Monitors

The combination of

- an object and its methods
- a mutual exclusion lock
- and the lock's condition objects

is called a **monitor**

Monitors enable modular synchronization.

Java's Lock Interface

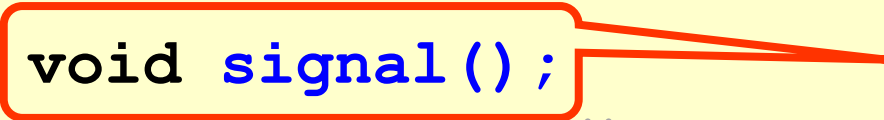
```
public interface Lock {  
    void lock();  
    void lockInterruptibly()  
        throws InterruptedException;  
    void tryLock();  
    void tryLock(long time, TimeUnit unit)  
        Condition newCondition();  
    void unlock();  
}
```

Java's Condition Interface

```
public interface Condition {  
    void await() throws InterruptedException;  
    boolean await(long time, TimeUnit unit)  
        throws InterruptedException;  
  
    ...  
    void signal();  
    void signalAll();  
}
```


Java's Condition Interface


```
public interface Condition {  
    void await() throws InterruptedException;  
    boolean await(long time, TimeUnit unit)  
        throws InterruptedException;  
  
    ...  
    void signal() ;  
    void signalAll() ;  
}
```



wake up **one**
waiting thread

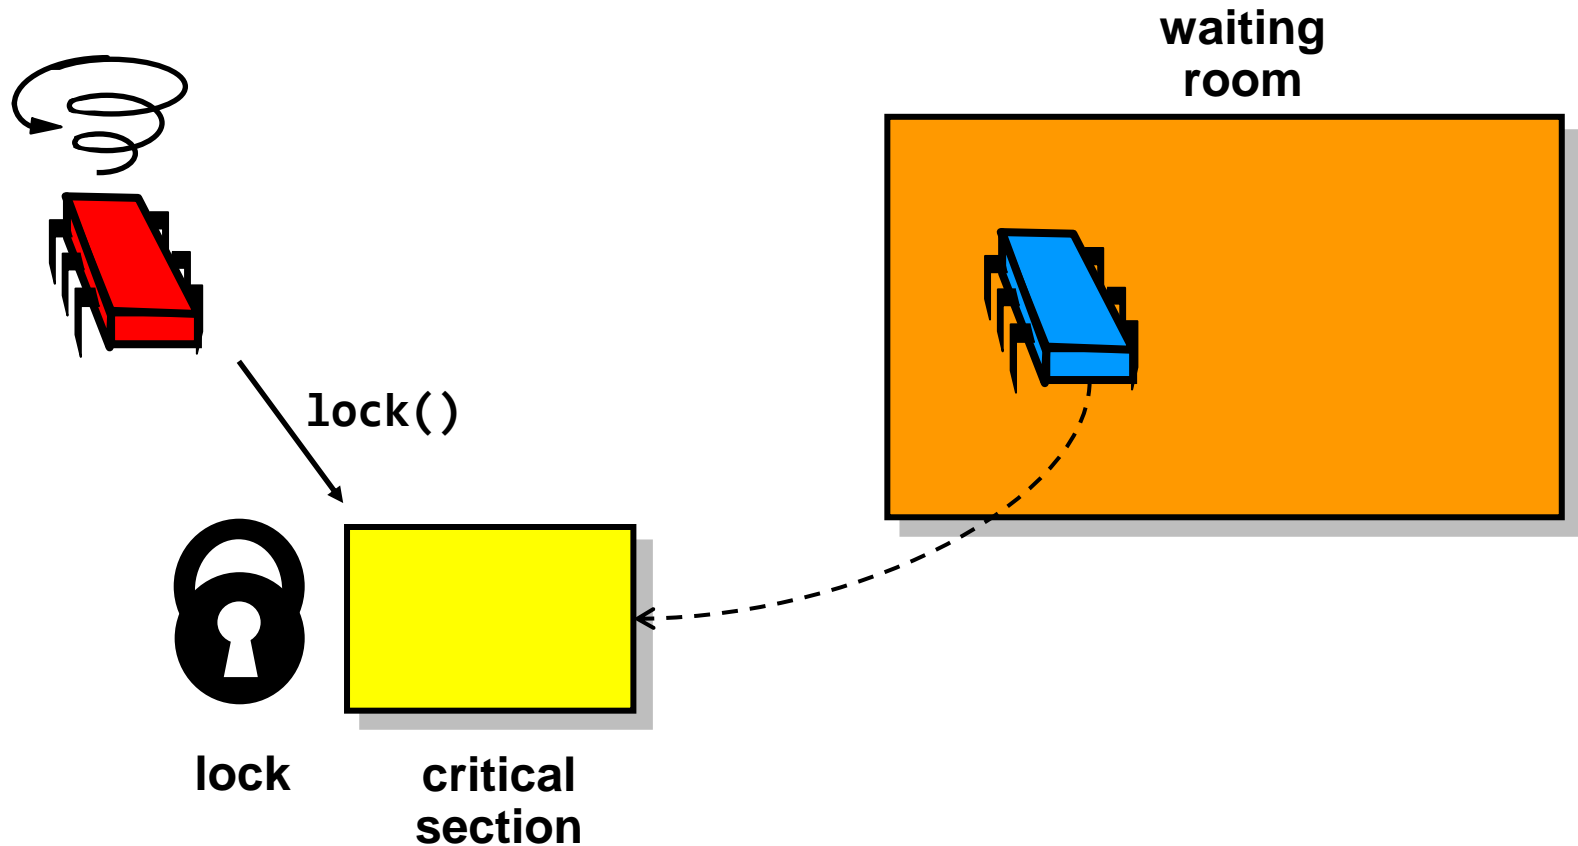
Java's Condition Interface

```
public interface Condition {  
    void await() throws InterruptedException;  
    boolean await(long time, TimeUnit unit)  
        throws InterruptedException;  
    ...  
    void signal();  
    void signalAll() ;  
}
```




wake up **all**
waiting threads

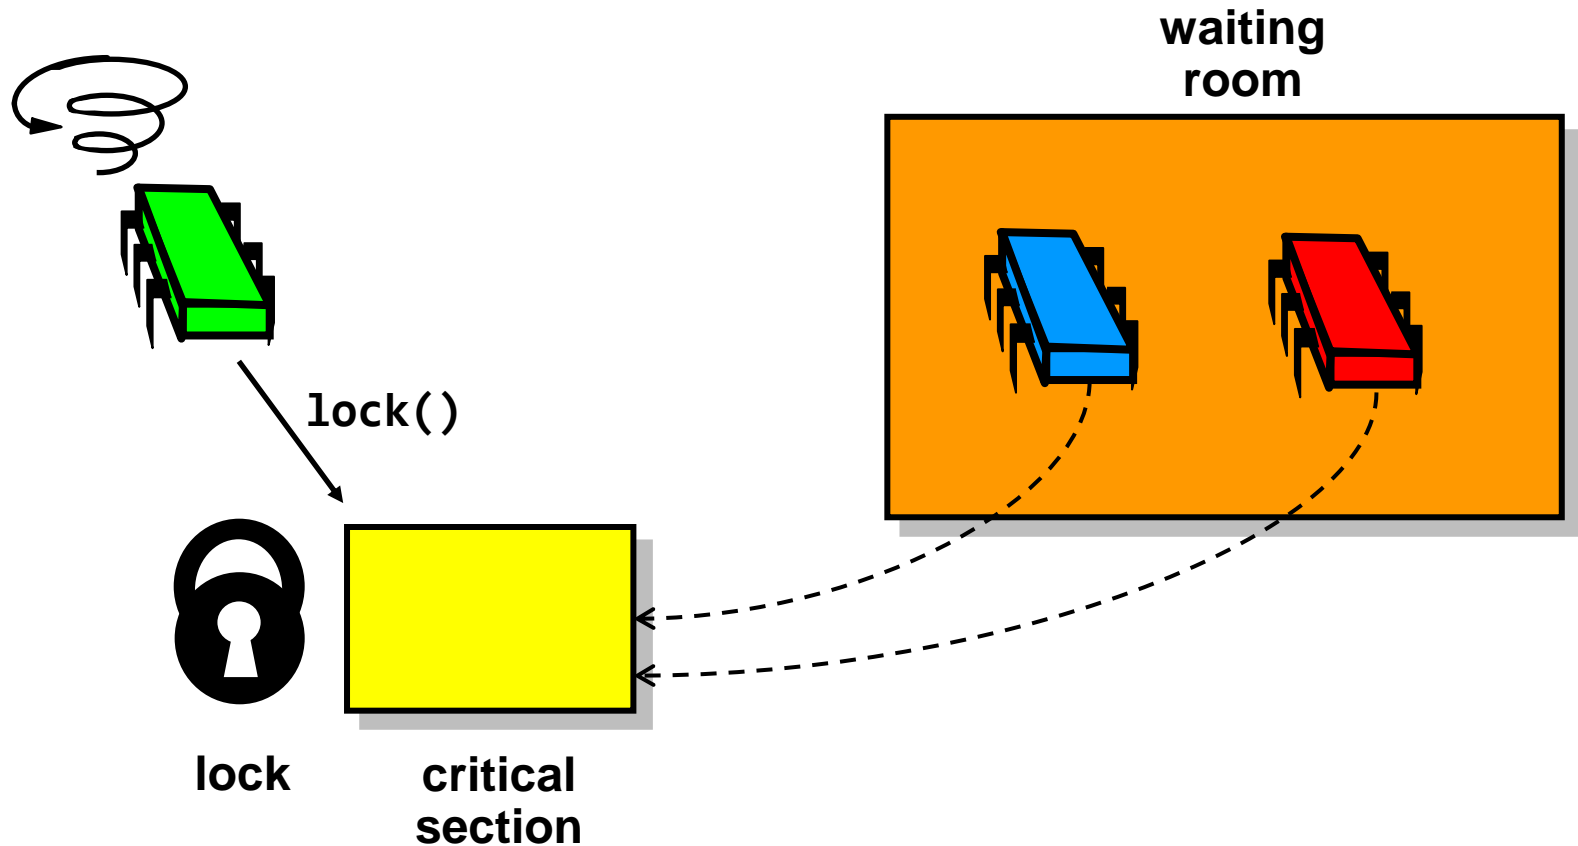
A Typical Monitor Execution




A Typical Monitor Execution




A Typical Monitor Execution




A Typical Monitor Execution




A Typical Monitor Execution




A Typical Monitor Execution



A Typical Monitor Execution



A Typical Monitor Execution



Using Condition Objects

```
Condition condition = mutex.newCondition();  
...  
mutex.lock();  
try {  
    while (!property)  
        condition.await();  
} catch (InterruptedException e) {  
    ...  
}  
...
```

Using Condition Objects

```
Condition condition = mutex.newCondition();
```

```
...
```

```
mutex.lock();
```

create new condition object

```
try {
```

```
    while (!property)
```

```
        condition.await();
```

```
} catch (InterruptedException e) {
```

```
    ...
```

```
}
```

```
...
```

Using Condition Objects

```
Condition condition = mutex.newCondition();
```

```
...
```

```
mutex.lock();
```

acquire the lock

```
try {
```

```
    while (!property)
```

```
        condition.await();
```

```
} catch (InterruptedException e) {
```

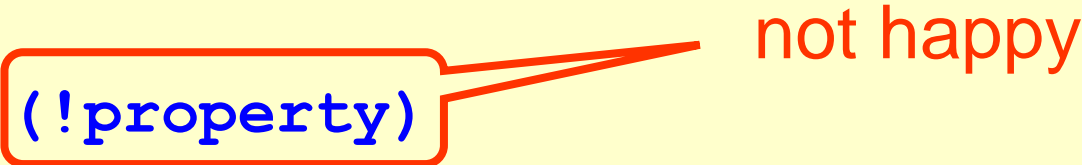
```
    ...
```

```
}
```

```
...
```

Using Condition Objects

```
Condition condition = mutex.newCondition();
...
mutex.lock();
try {
    while (!property)
        condition.await();
} catch (InterruptedException e) {
    ...
}
...
```



not happy


Using Condition Objects

```
Condition condition = mutex.newCondition();  
...  
mutex.lock();  
try {  
    while (!property)  
        condition.await();  
} catch (InterruptedException e) {  
    ...  
}  
...
```

release the lock
and suspend
until notified

Using Condition Objects

```
Condition condition = mutex.newCondition();  
...  
mutex.lock();  
try {  
    while (!property)  
        condition.await();  
} catch (InterruptedException e) {  
    ...  
}  
...
```



application specific response

Using Condition Objects

```
Condition condition = mutex.newCondition();  
...  
mutex.lock();  
try {  
    while (!property)  
        condition.await();  
} catch (InterruptedException e) {  
    ...  
}
```

...

happy: **property** must hold

Example: Blocking Queue

```
public class BlockingQueue<T> {  
    final Lock lock = new ReentrantLock();  
    final Condition notFull = lock.newCondition();  
    final Condition notEmpty = lock.newCondition();  
    final T[] items;  
    int tail, head, count;  
  
    public BlockingQueue(int capacity) {  
        items = new T[capacity];  
    }  
    ...  
}
```

Example: Blocking Queue

```
public class BlockingQueue<T> {  
    final Lock lock = new ReentrantLock();  
    final Condition notFull = lock.newCondition();  
    final Condition notEmpty = lock.newCondition();  
    final T[] items;  
    int tail, head, count;  
  
    public BlockingQueue(int capacity) {  
        items = new T[capacity];  
    }  
    ...  
}
```

mutual exclusion lock
for queue object

Example: Blocking Queue

```
public class BlockingQueue<T> {  
    final Lock lock = new ReentrantLock();  
    final Condition notFull = lock.newCondition();  
    final Condition notEmpty = lock.newCondition();  
    final T[] items;  
    int tail, head, count;  
  
    public BlockingQueue(int capacity) {  
        items = new T[capacity];  
    }  
    ...  
}
```

condition to wait on
if queue is full

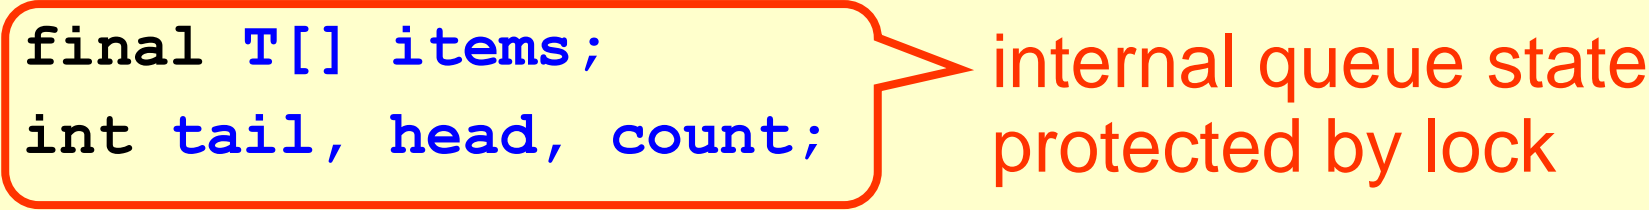
Example: Blocking Queue

```
public class BlockingQueue<T> {  
    final Lock lock = new ReentrantLock();  
    final Condition notFull = lock.newCondition();  
    final Condition notEmpty = lock.newCondition();  
    final T[] items;  
    int tail, head, count;  
  
    public BlockingQueue(int capacity) {  
        items = new T[capacity];  
    }  
    ...  
}
```

condition to wait on
if queue is empty

Example: Blocking Queue

```
public class BlockingQueue<T> {  
    final Lock lock = new ReentrantLock();  
    final Condition notFull = lock.newCondition();  
    final Condition notEmpty = lock.newCondition();  
    final T[] items;  
    int tail, head, count;  
  
    public BlockingQueue(int capacity) {  
        items = new T[capacity];  
    }  
    ...  
}
```



internal queue state
protected by lock

Example: Blocking Queue

```
public class BlockingQueue<T> {  
    final Lock lock = new ReentrantLock();  
    final Condition notFull = lock.newCondition();  
    final Condition notEmpty = lock.newCondition();  
    final AtomicReferenceArray<T> items;  
    volatile int tail, head, count;  
  
    public BlockingQueue(int capacity) {  
        items = new T[capacity];  
    }  
    ...  
}
```

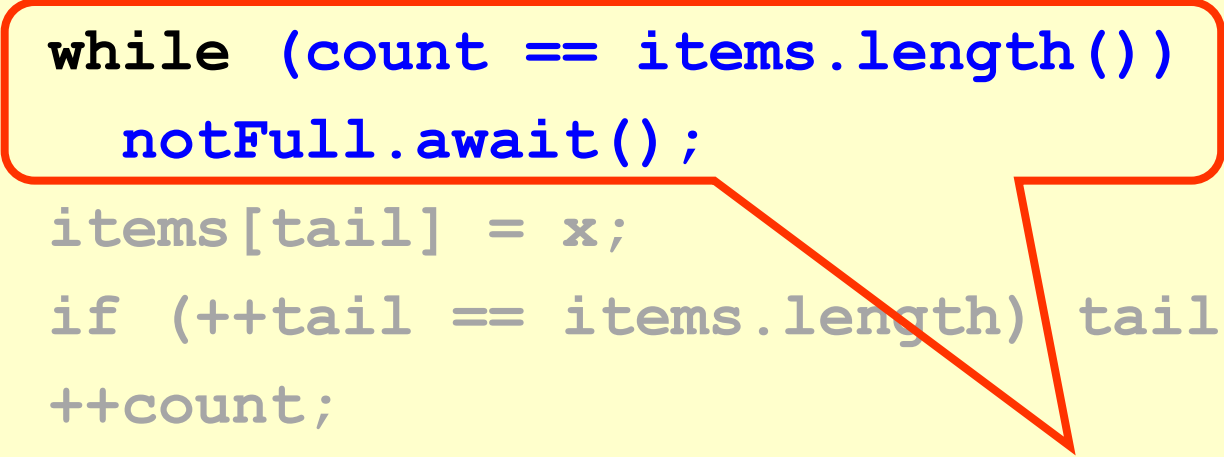
beware of weak memory consistency due to caching

Blocking Queue: enqueue

```
public void enq(T x) {  
    lock.lock();  
    try {  
        while (count == items.length())  
            notFull.await();  
        items[tail] = x;  
        if (++tail == items.length) tail = 0;  
        ++count;  
        notEmpty.signal();  
    } finally { lock.unlock(); }  
}
```


Blocking Queue: enqueue

```
public void enq(T x) {  
    lock.lock();  
    try {  
        while (count == items.length())  
            notFull.await();  
        items[tail] = x;  
        if (++tail == items.length) tail = 0;  
        ++count;  
        notEmpty.signal();  
    } finally { lock.unlock(); }  
}
```




wait until queue
has space

Blocking Queue: enqueue

```
public void enq(T x) {  
    lock.lock();  
    try {  
        while (count == items.length())  
            notFull.await();  
  
        items[tail] = x;  
        if (++tail == items.length) tail = 0;  
        ++count;  
  
        notEmpty.signal();  
    } finally { lock.unlock(); }  
}
```

queue has space!
insert element

Blocking Queue: enqueue

```
public void enq(T x) {  
    lock.lock();  
    try {  
        while (count == items.length())  
            notFull.await();  
        items[tail] = x;  
        if (++tail == items.length) tail = 0;  
        ++count;  
        notEmpty.signal();  wake up one waiting  
                                consumer  
    } finally { lock.unlock(); }  
}
```

Blocking Queue: dequeue

```
public T deq() {  
    lock.lock();  
    try {  
        while (count == 0)  
            notEmpty.await();  
        T x = items[head];  
        if (++head == items.length) head = 0;  
        --count;  
        notFull.signal();  
        return x;  
    } finally { lock.unlock(); }  
}
```

Blocking Queue: dequeue

```
public T deq() {  
    lock.lock();  
    try {  
        while (count == 0)  
            notEmpty.await();  
        T x = items[head];  
        if (++head == items.length) head = 0;  
        --count;  
        notFull.signal();  
        return x;  
    } finally { lock.unlock(); }  
}
```

wait until queue
is nonempty

Blocking Queue: dequeue

```
public T deq() {  
    lock.lock();  
    try {  
        while (count == 0)  
            notEmpty.await();  
  
        T x = items[head];  
        if (++head == items.length) head = 0;  
        --count;  
  
        notFull.signal();  
        return x;  
    } finally { lock.unlock(); }  
}
```

queue nonempty!
retrieve next
element

Blocking Queue: dequeue

```
public T deq() {  
    lock.lock();  
    try {  
        while (count == 0)  
            notEmpty.await();  
        T x = items[head];  
        if (++head == items.length) head = 0;  
        --count;  
        notFull.signal();  
        return x;  
    } finally { lock.unlock(); }  
}
```

wake up one waiting producer

Improved enqueue?


```
public void enq(T x) {  
    lock.lock();  
    try {  
        while (count == items.length())  
            notFull.await();  
        items[tail] = x;  
        if (++tail == items.length) tail = 0;  
        ++count;  
        if (count == 1) notEmpty.signal();  
    } finally { lock.unlock(); }  
}
```


Improved enqueue?


```
public void enq(T x) {  
    lock.lock();  
    try {  
        while (count == items.length())  
            notFull.await();  
        items[tail] = x;  
        if (++tail == items.length) tail = 0;  
        ++count;  
        if (count == 1) notEmpty.signal();  
    } finally { lock.unlock(); }  
}
```

lost wakeups


Lost Wakeup




Lost Wakeup




Lost Wakeup




Lost Wakeup




Lost Wakeup




Lost Wakeup




Lost Wakeup




Lost Wakeup




Lost Wakeup




Lost Wakeup




Lost Wakeup



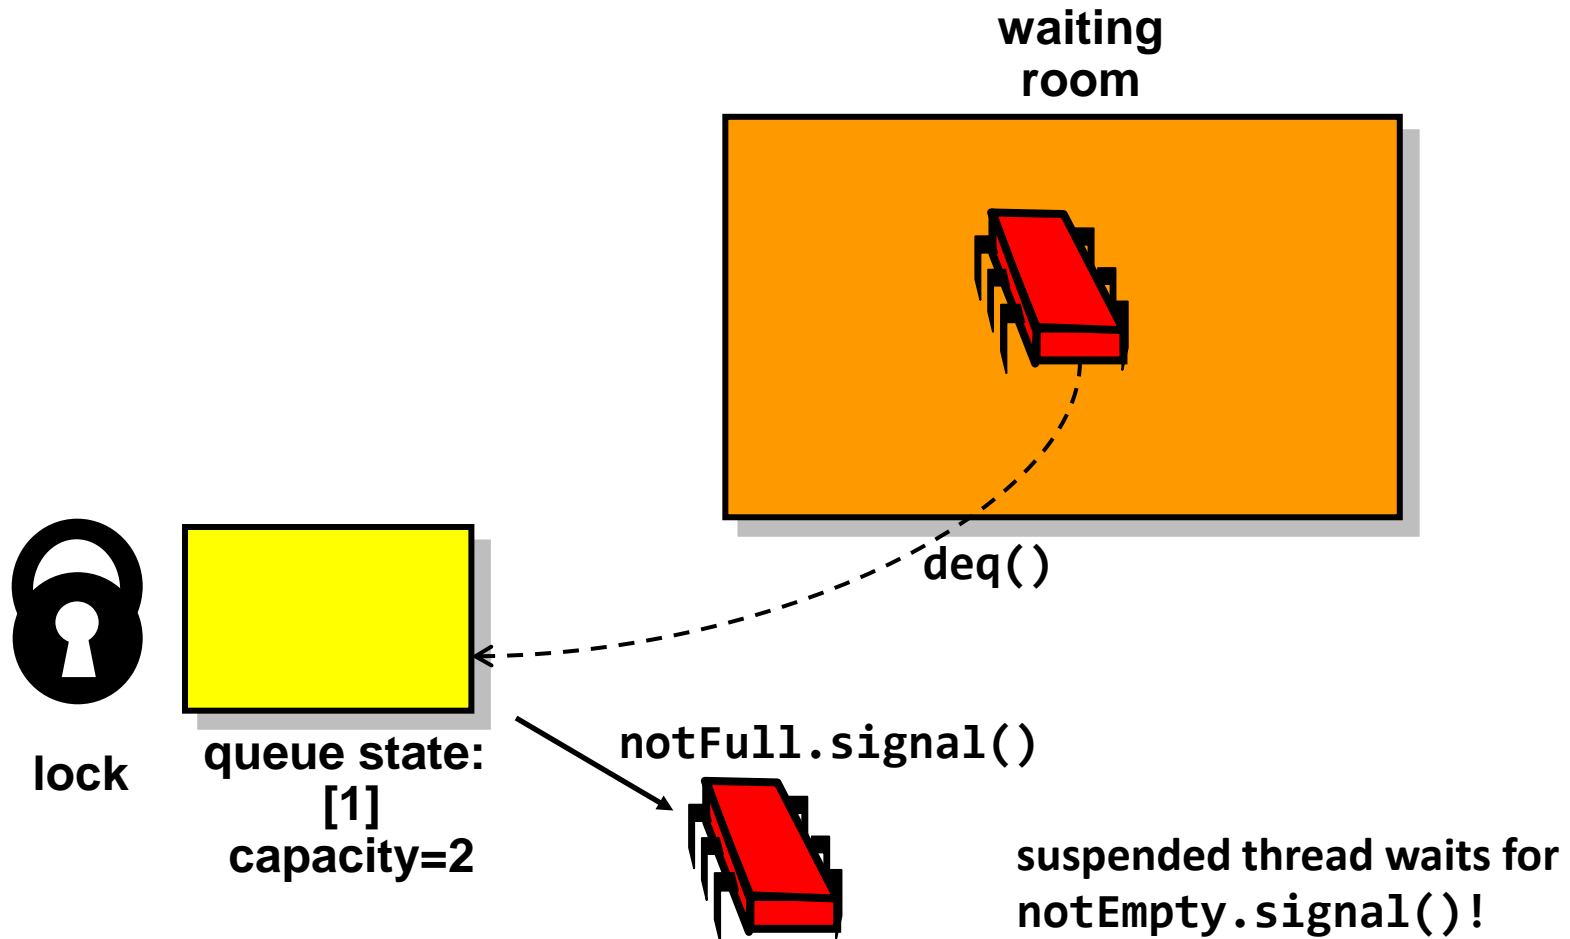
Lost Wakeup




Lost Wakeup



Lost Wakeup



Lost Wakeup



The Lost-Wakeup Problem

- Condition variables are inherently vulnerable to lost wakeups
 - one thread waits forever without realizing that its waiting condition has become true
- Programming practices
 - if in doubt, signal **all** waiting processes
 - specify a timeout when waiting

Reentrant Locks

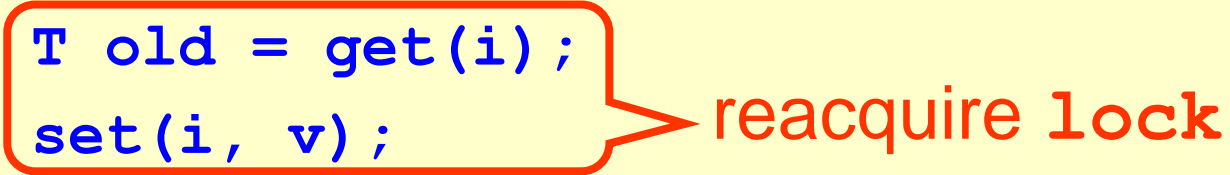
- same thread can acquire the lock multiple times without blocking
- commonly used in OOP to handle reentrant calls to locked objects

Using Reentrant Locks

```
public class AtomicArray<T> {  
    final Lock lock = new ReentrantLock();  
    ...  
    public T getAndSet(int i, T v) {  
        try { lock.lock();  
            T old = get(i);  
            set(i, v);  
            return old;  
        } finally { lock.unlock(); } }  
    public T get() {  
        try {lock.lock(); return item[i]; }  
        finally { lock.unlock(); }  
    }  
    public void set(int i, T v) { ... } }  
}
```

Using Reentrant Locks

```
public class AtomicArray<T> {  
    final Lock lock = new ReentrantLock();  
    ...  
    public T getAndSet(int i, T v) {  
        try { lock.lock();  
            T old = get(i);  
            set(i, v);  
            return old;  
        } finally { lock.unlock(); } }  
    public T get() {  
        try {lock.lock(); return item[i]; }  
        finally { lock.unlock(); }  
    }  
    public void set(int i, T v) { ... } }  
}
```



reacquire lock

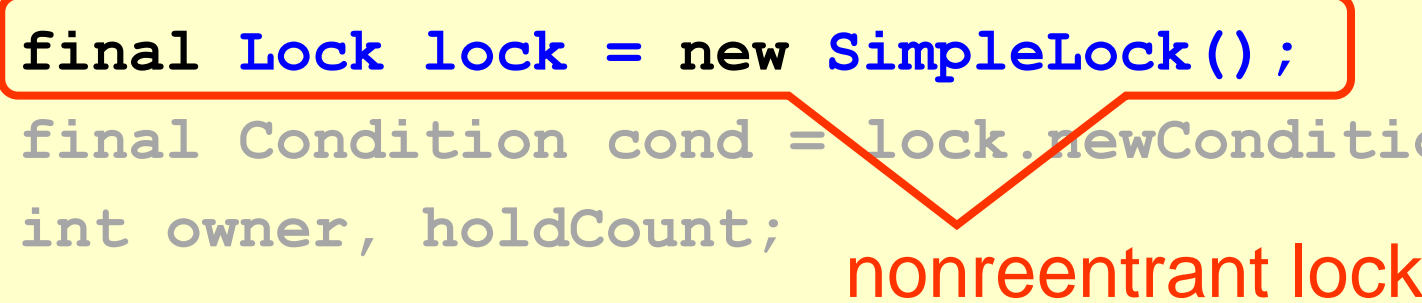
Our Own Reentrant Lock

```
public class SimpleReentrantLock implements Lock{
    final Lock lock = new SimpleLock();
    final Condition cond = lock.newCondition();
    int owner, holdCount;

    public SimpleReentrantLock() {
        owner = holdCount = 0;
    }
    ...
}
```

Our Own Reentrant Lock

```
public class SimpleReentrantLock implements Lock{  
    final Lock lock = new SimpleLock();  
    final Condition cond = lock.newCondition();  
    int owner, holdCount;  
  
    public SimpleReentrantLock() {  
        owner = holdCount = 0;  
    }  
    ...  
}
```

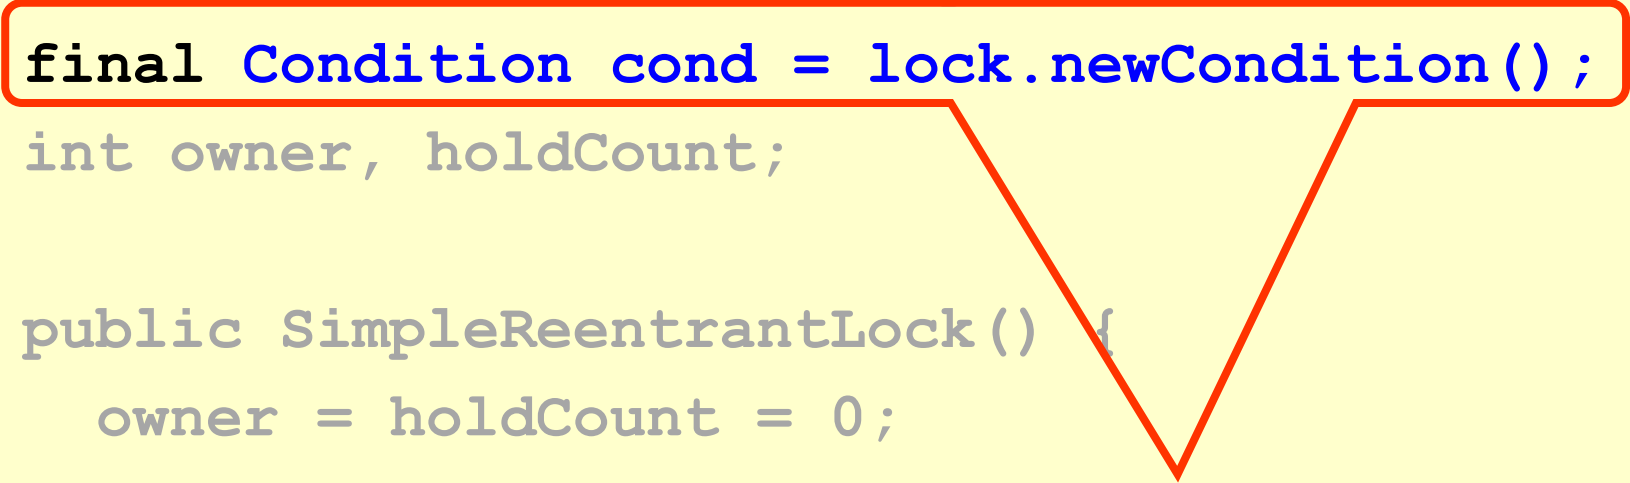


nonreentrant lock

Our Own Reentrant Lock

```
public class SimpleReentrantLock implements Lock{
    final Lock lock = new SimpleLock();
    final Condition cond = lock.newCondition();
    int owner, holdCount;

    public SimpleReentrantLock() {
        owner = holdCount = 0;
    }
    ...
}
```

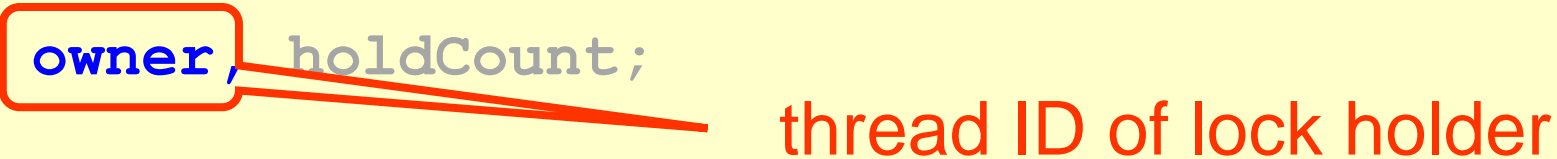


condition to wait on if lock
is held by other thread

Our Own Reentrant Lock

```
public class SimpleReentrantLock implements Lock{
    final Lock lock = new SimpleLock();
    final Condition cond = lock.newCondition();
    int owner, holdCount;

    public SimpleReentrantLock() {
        owner = holdCount = 0;
    }
    ...
}
```



thread ID of lock holder

Our Own Reentrant Lock

```
public class SimpleReentrantLock implements Lock{
    final Lock lock = new SimpleLock();
    final Condition cond = lock.newCondition();
    int owner, holdCount;
    public SimpleReentrantLock() {
        owner = holdCount = 0;
    }
    ...
}
```

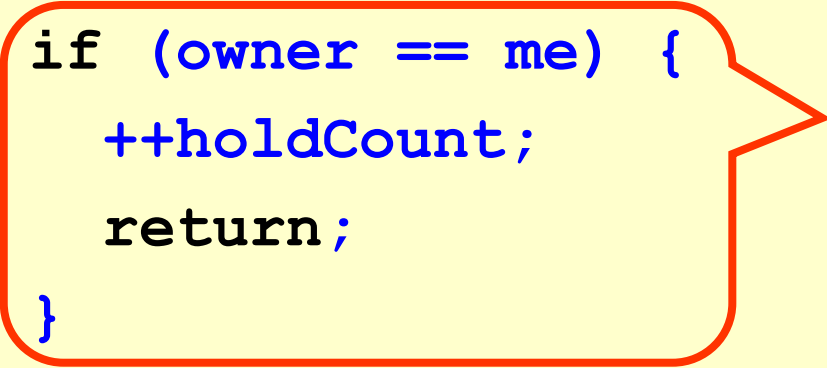
counts how often lock
has been acquired by
current owner

Our Own Reentrant Lock

```
public void lock() {  
    int me = ThreadID.get();  
    lock.lock();  
    try {  
        if (owner == me) {  
            ++holdCount;  
            return;  
        }  
        while (holdCount != 0) condition.await();  
        owner = me;  
        holdCount = 1;  
    } finally { lock.unlock() } }
```

Our Own Reentrant Lock

```
public void lock() {  
    int me = ThreadID.get();  
    lock.lock();  
    try {  
        if (owner == me) {  
            ++holdCount;  
            return;  
        }  
        while (holdCount != 0) condition.await();  
        owner = me;  
        holdCount = 1;  
    } finally { lock.unlock() } }
```



already holding the lock?
then just increase counter

Our Own Reentrant Lock

```
public void lock() {  
    int me = ThreadID.get();  
    lock.lock();  
    try {  
        if (owner == me) {  
            ++holdCount;           otherwise, wait until lock is  
            return;                free and then take ownership  
        }  
  
        while (holdCount != 0) condition.await();  
        owner = me;  
        holdCount = 1;  
    } finally { lock.unlock() } }
```

Our Own Reentrant Lock

```
public void unlock() {  
    lock.lock();  
    try {  
        if (holdCount == 0 ||  
            owner != ThreadID.get()) {  
            throw new IllegalMonitorStateException();  
        }  
        if (--holdCount == 0) cond.signal();  
    } finally { lock.unlock(); }  
}
```

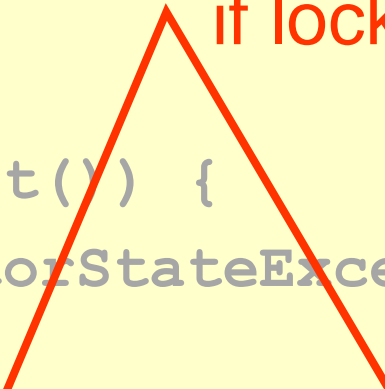
Our Own Reentrant Lock

```
public void unlock() {  
    lock.lock();          fail, if lock is released too often  
    try {  
        if (holdCount == 0 ||  
            owner != ThreadID.get()) {  
            throw new IllegalMonitorStateException();  
        }  
        if (--holdCount == 0) cond.signal();  
    } finally { lock.unlock() }  
}
```

Our Own Reentrant Lock

```
public void unlock() {  
    lock.lock();  
    try {  
        if (holdCount == 0 ||  
            owner != ThreadID.get()) {  
            throw new IllegalMonitorStateException();  
        }  
        if (--holdCount == 0) cond.signal();  
    } finally { lock.unlock(); }  
}
```

otherwise, decrement counter
and wake up one blocked thread
if lock is released



Java's built-in Monitors

- synchronized blocks, and methods acquire and release an implicit reentrant lock
- access to an implicit condition object is provided via special methods
 - wait()
 - notify()
 - notifyAll()

Simplified Blocking Queue: enqueue


```
public synchronized void enq(T x) {  
    while (count == items.length())  
        wait();  
    items[tail] = x;  
    if (++tail == items.length) tail = 0;  
    ++count;  
    notifyAll();  
}
```

Simplified Blocking Queue: dequeue

```
public synchronized T deq() {  
    while (count == 0)  
        wait();  
    T x = items[head];  
    if (++head == items.length) head = 0;  
    --count;  
    notifyAll();  
    return x;  
}
```

Simplified Blocking Queue: dequeue

```
public synchronized T deq() {  
    while (count == 0)  
        wait();  
    T x = items[head];  
    if (++head == items.length) head = 0;  
    --count;  
    notify();  
    return x;  
}
```



is notify enough?


Simplified Blocking Queue: dequeue

```
public synchronized T deq() {  
    while (count == 0)  
        wait();  
    T x = items[head];  
    if (++head == items.length) head = 0;  
    --count;  
    notify();  
    return x;  
}
```


is `notify` enough?

lost wakeups

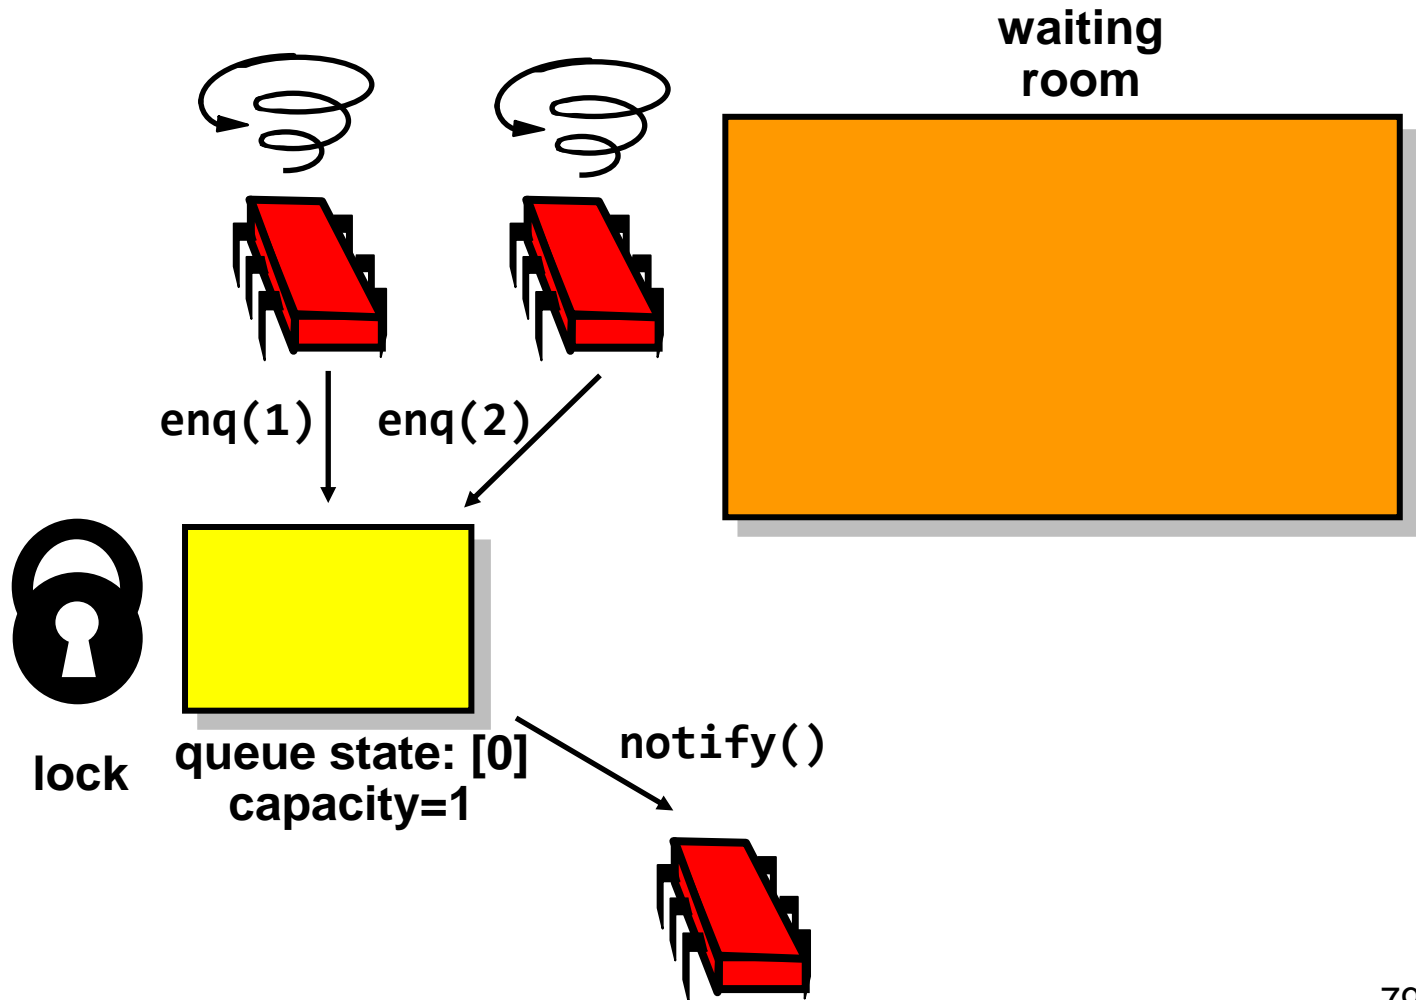
Lost Wakeup in Simplified Queue with `notify()`




Lost Wakeup in Simplified Queue with `notify()`




Lost Wakeup in Simplified Queue with `notify()`




Lost Wakeup in Simplified Queue with `notify()`



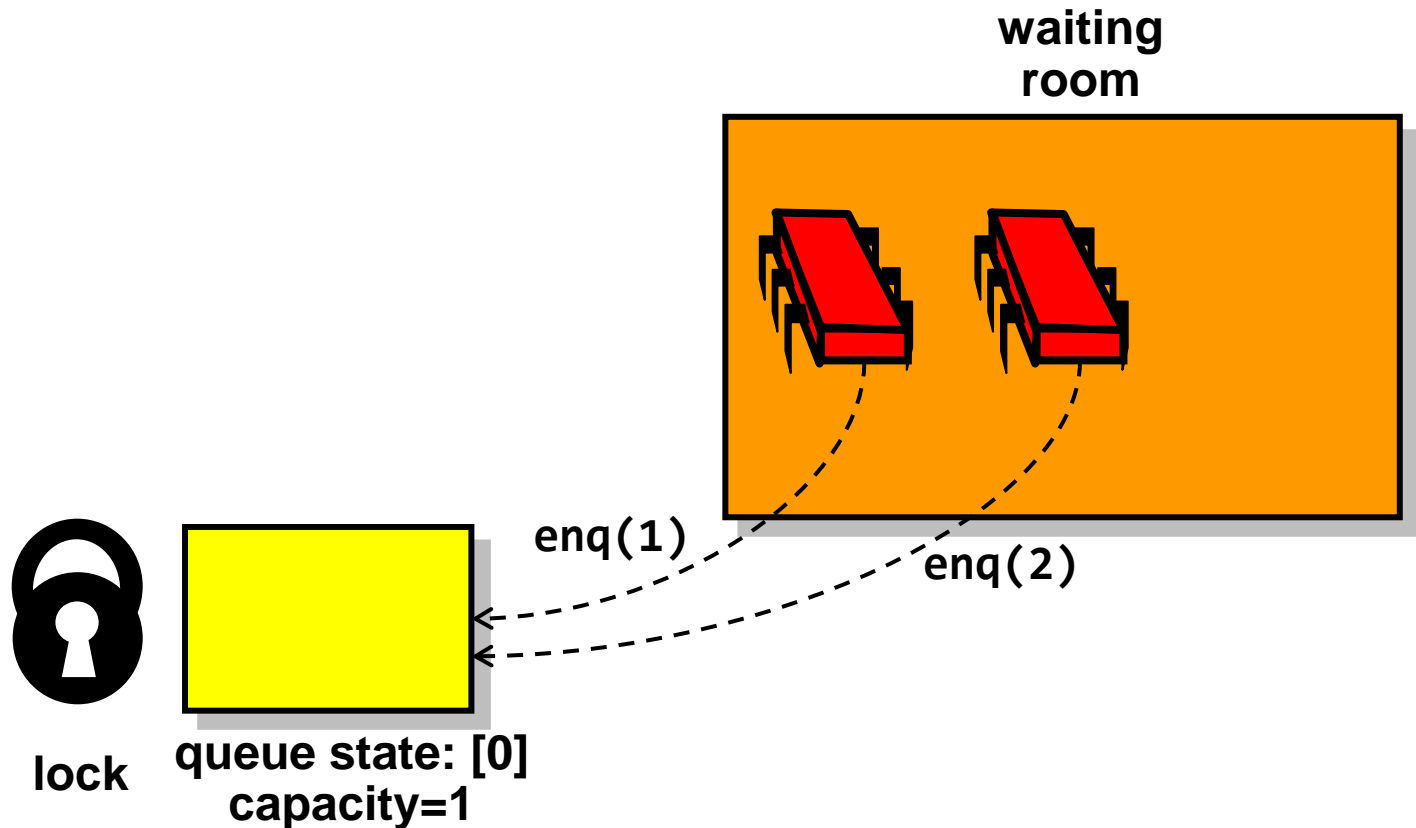
Lost Wakeup in Simplified Queue with `notify()`




Lost Wakeup in Simplified Queue with `notify()`




Lost Wakeup in Simplified Queue with `notify()`




Lost Wakeup in Simplified Queue with notify()




Lost Wakeup in Simplified Queue with notify()




Lost Wakeup in Simplified Queue with `notify()`




Lost Wakeup in Simplified Queue with notify()




Lost Wakeup in Simplified Queue with `notify()`




Lost Wakeup in Simplified Queue with `notify()`



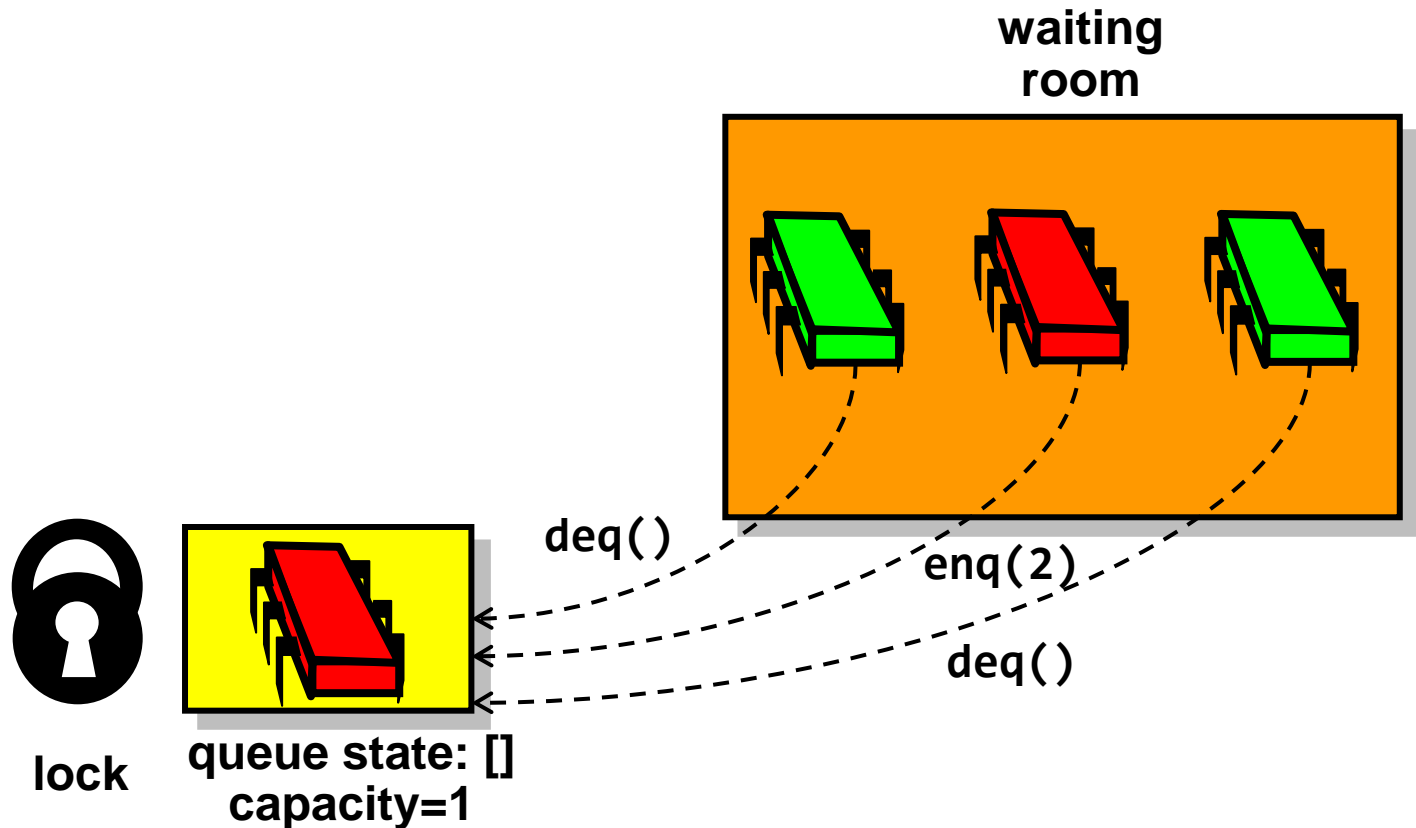
Lost Wakeup in Simplified Queue with notify()



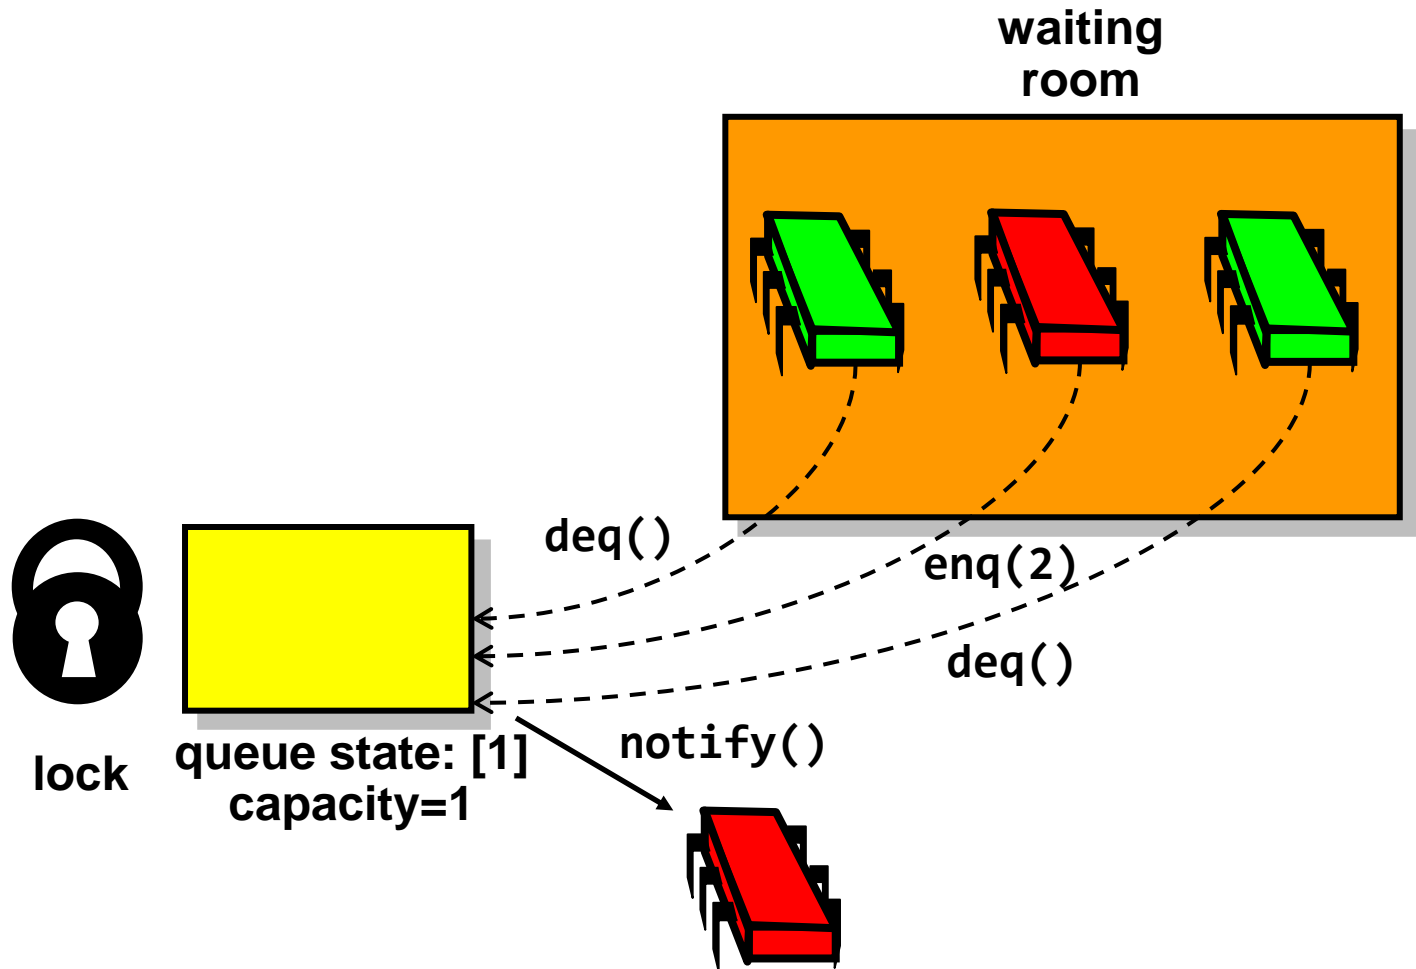
Lost Wakeup in Simplified Queue with notify()



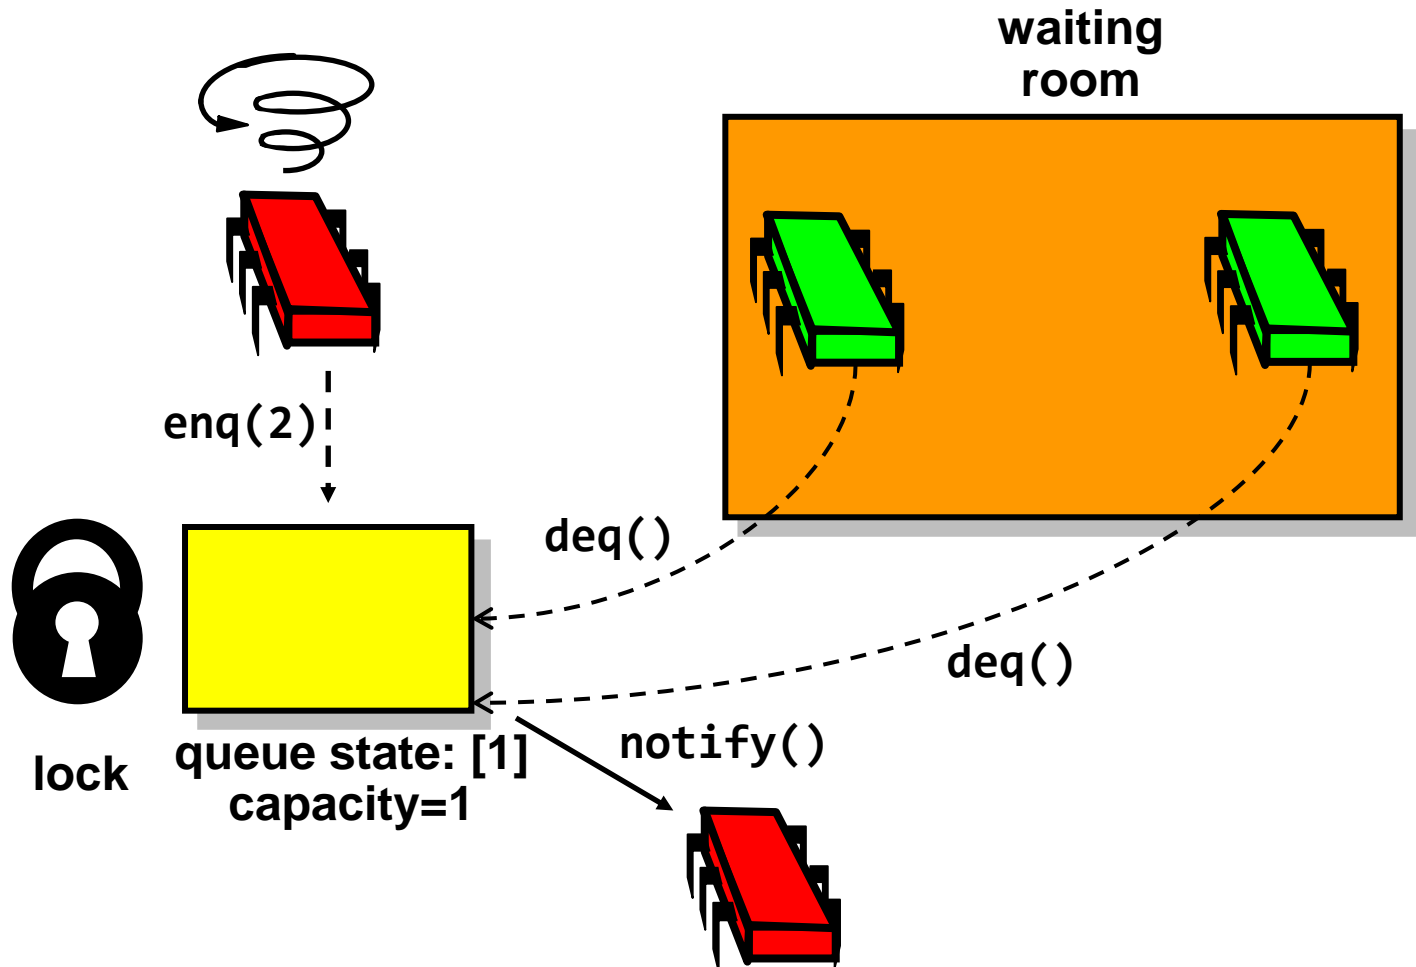
Lost Wakeup in Simplified Queue with notify()



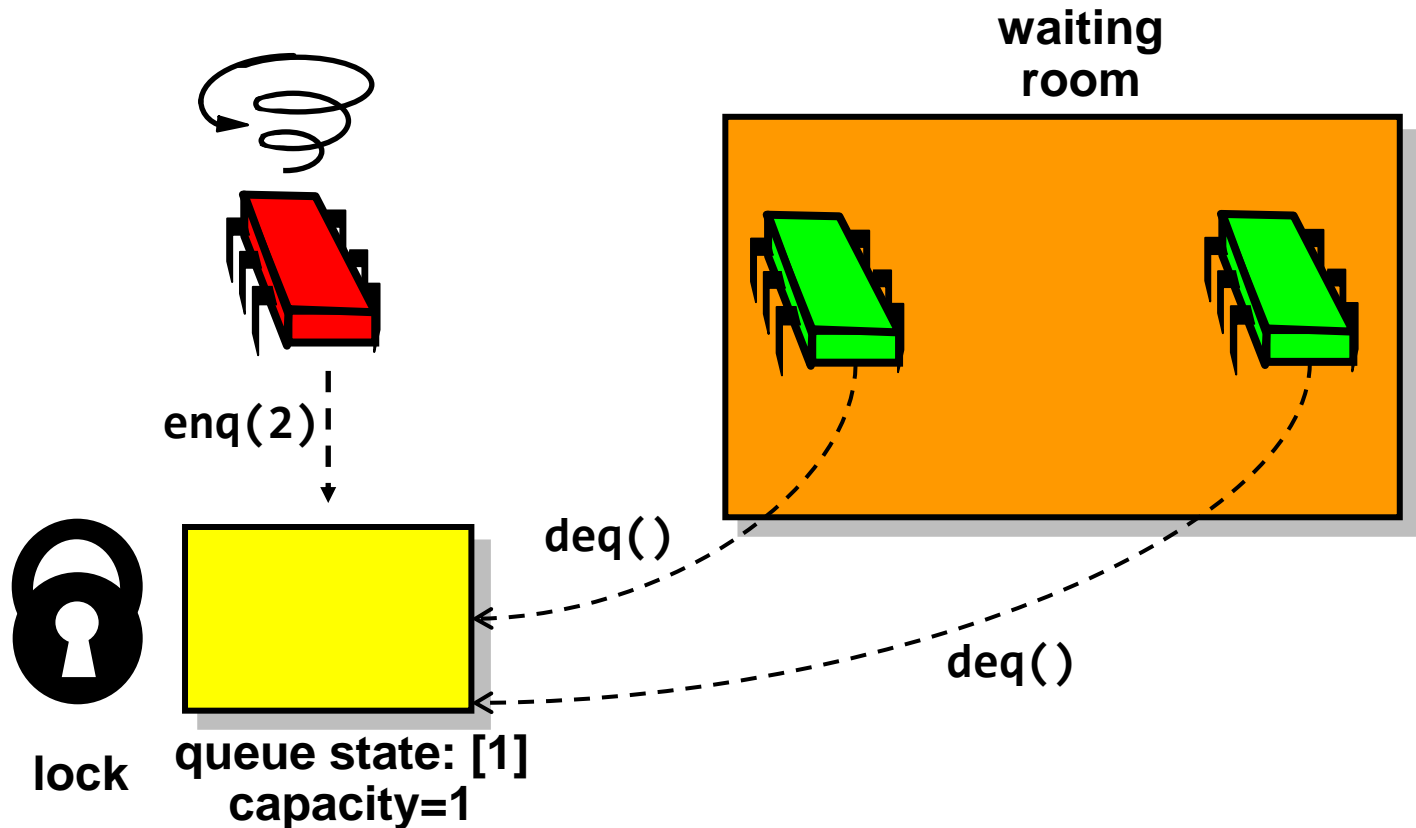
Lost Wakeup in Simplified Queue with notify()




Lost Wakeup in Simplified Queue with `notify()`



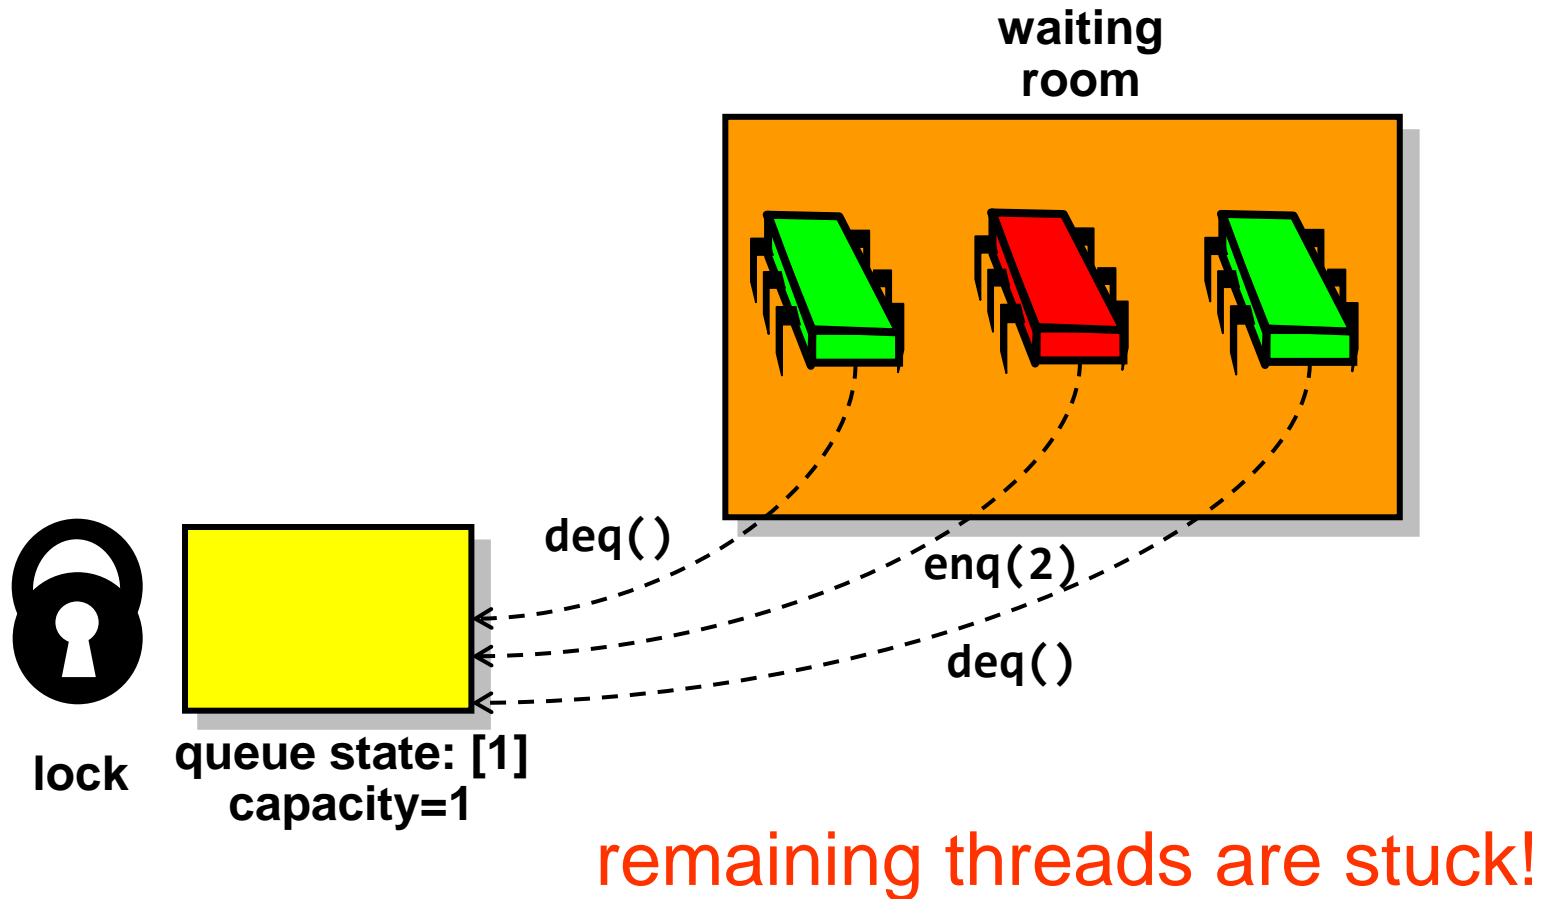
Lost Wakeup in Simplified Queue with `notify()`



Lost Wakeup in Simplified Queue with `notify()`



Lost Wakeup in Simplified Queue with `notify()`



Issues with Shared Memory Concurrency

- Complicated locking protocols
 - deadlocks
 - livelocks
 - lost wake-up
 - ...
- Weak memory consistency guarantees: need to use
 - `volatile` variables
 - `AtomicReferenceArray<T>` instead of `T[]`
 - ...

`java.lang.concurrent` helps but is not a panacea.